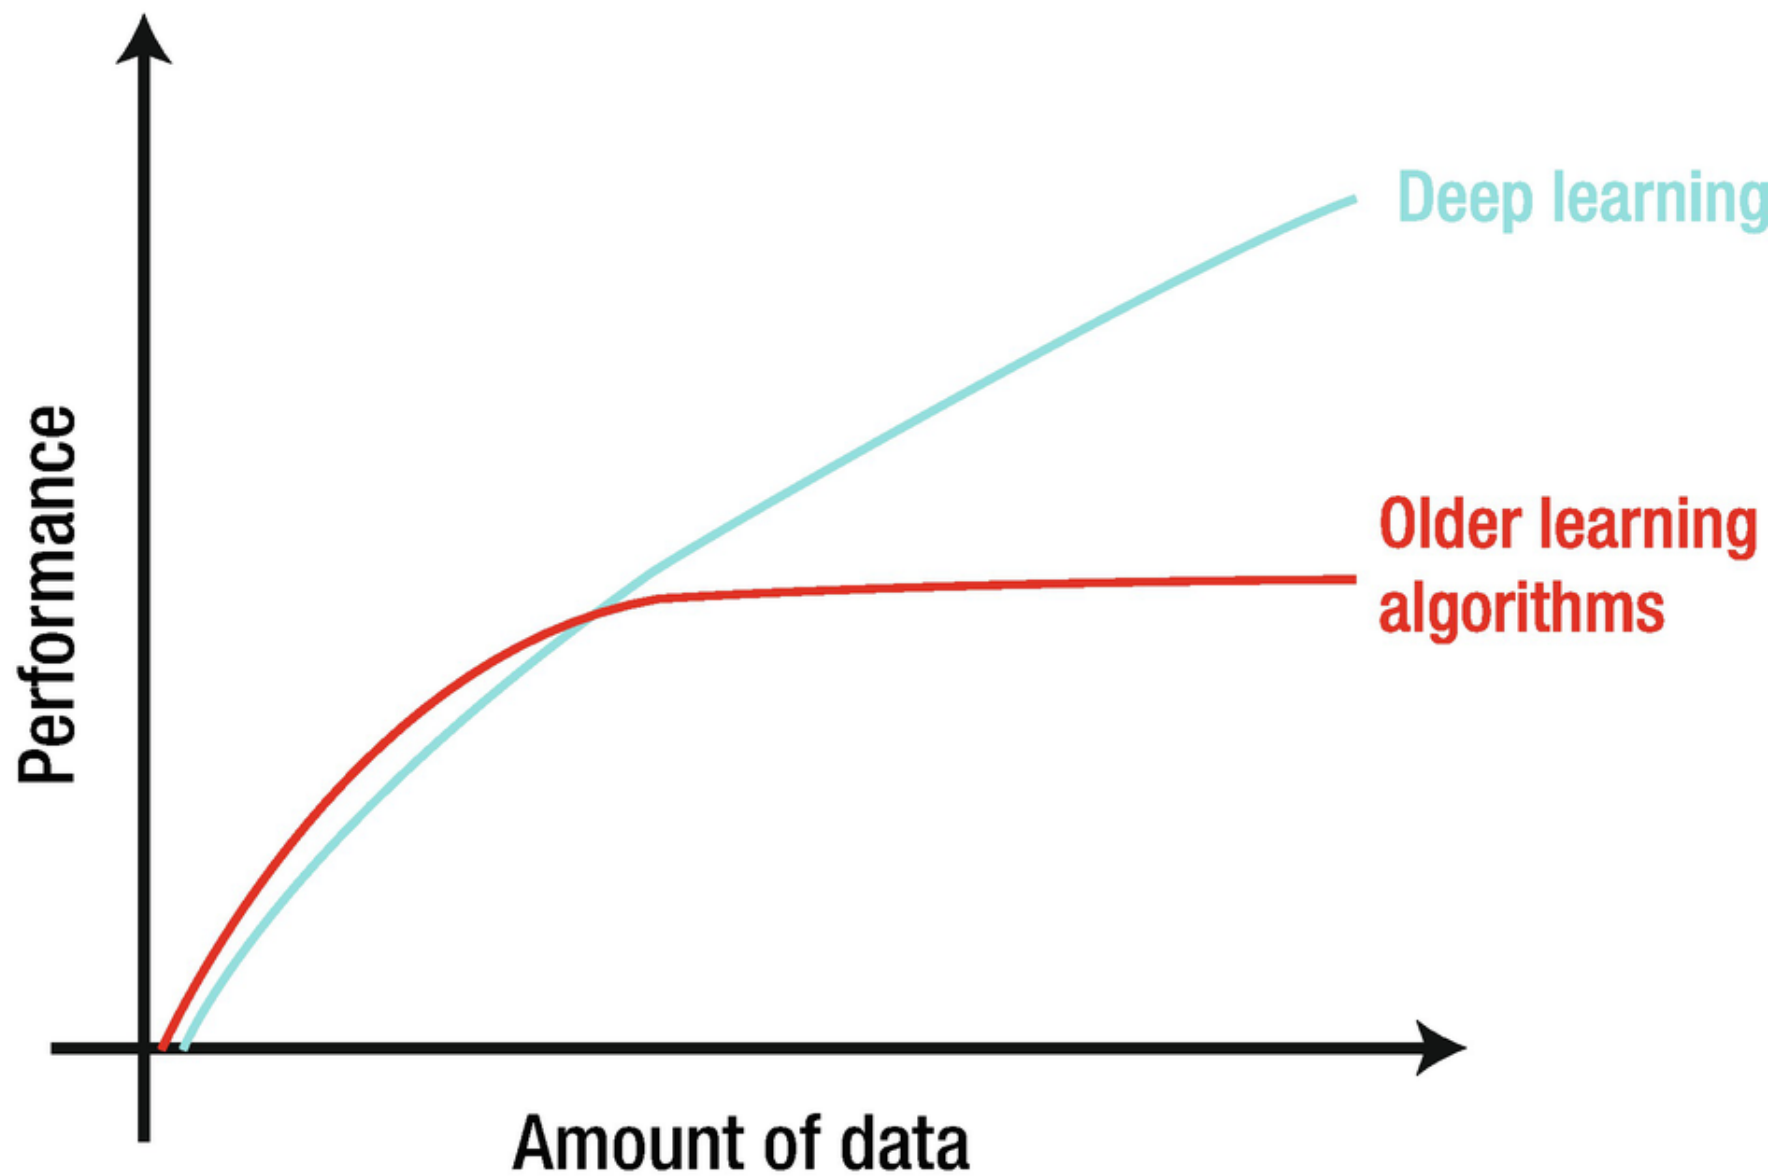


Improving Deep Neural Networks

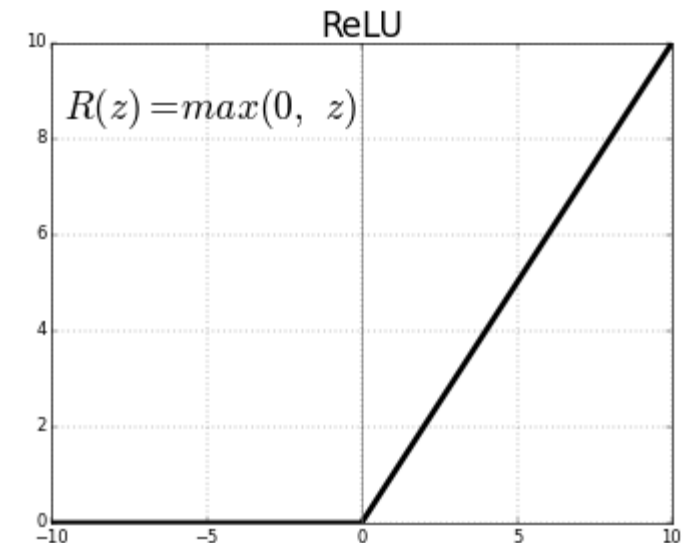
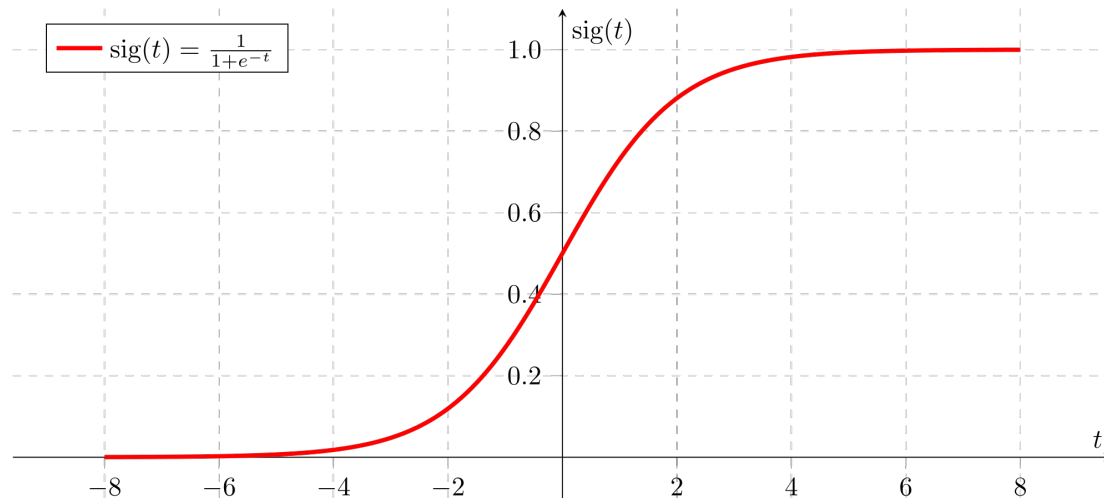
Il-Youp Kwak, PhD

Why deep learning?



Why deep learning taking off?

- Firstly proposed in 1943
- Originally had problem in computational speed
- Development in hardware for computing (GPUs) and algorithm itself



Binary Classification

- We are Classifying Cat or Dog $f : \mathbf{x} \xrightarrow{f_\theta} \mathbb{R}_{[0,1]}$
- Dimension for \mathbf{x} is $64*64*3 = 12288$
- Data: (\mathbf{x}, y) $\mathbf{x} \in \mathbb{R}^{n_x}, y \in \{0, 1\}$
 m training examples $\{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(m)}, y^{(m)})\}$
 $X = [\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}]$ $Y = [y^{(1)}, \dots, y^{(m)}]$

Logistic Regression

- Given $\mathbf{x} \in \mathbb{R}^{n_x}$, want $\hat{y} = P(y = 1|\mathbf{x}) \in \mathbb{R}^{[0,1]}$
- Parameters: $\mathbf{w} \in \mathbb{R}^{n_x}$, $b \in \mathbb{R}$
- Output: $\hat{y} = \sigma(\mathbf{w}^t \mathbf{x} + b)$, where $\sigma(z) = 1/(1 + e^{-z})$

Loss function

- **Squared error loss:** $L(\hat{y}, y) = (\hat{y} - y)^2$

- **Cross entropy loss:**

$$L(\hat{y}, y) = -(y \log(\hat{y}) + (1 - y) \log(1 - \hat{y}))$$

- **Cost for Logistic regression: Use Cross entropy loss**

$$C(W, b) = \sum_{i=1}^m L(\hat{y}_i, y_i)$$

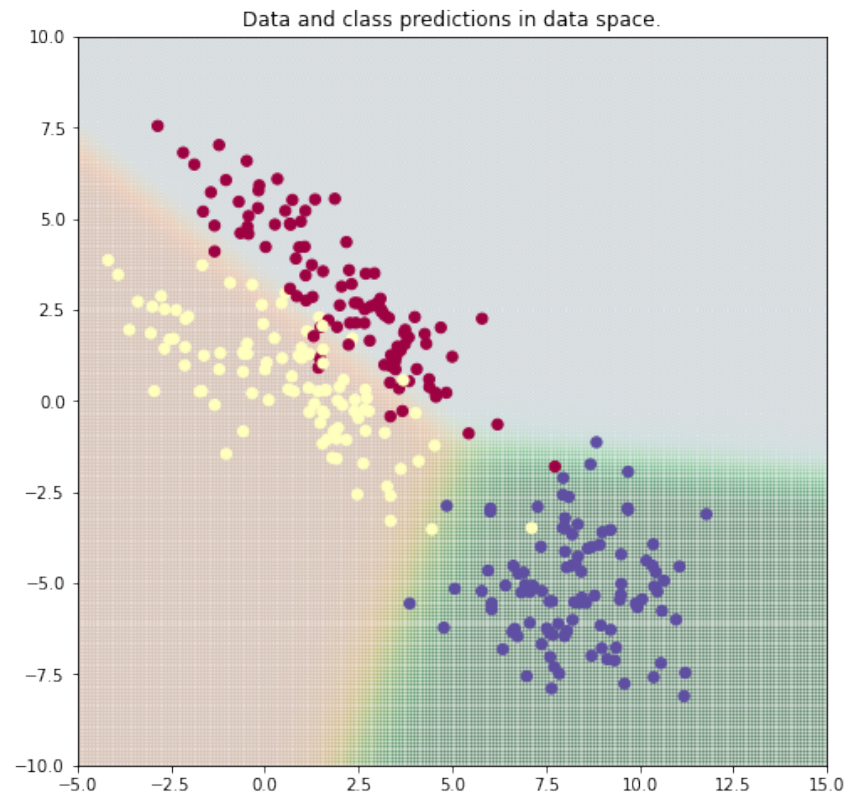
Constructing Logistic regression

```
model = keras.Sequential([  
    keras.layers.Flatten(input_shape=(28, 28)),  
    keras.layers.Dense(2, activation='sigmoid')  
])
```

```
[17] model.compile(optimizer='adam',  
                  loss='binary_crossentropy',  
                  metrics=['accuracy'])
```

Softmax Regression

- $\hat{\mathbf{z}} = e^{(W\mathbf{x}+\mathbf{b})}$ $t = \sum_i \hat{z}_i$
- Then, $\hat{y} = \hat{\mathbf{z}}/t$ represent probability for each item



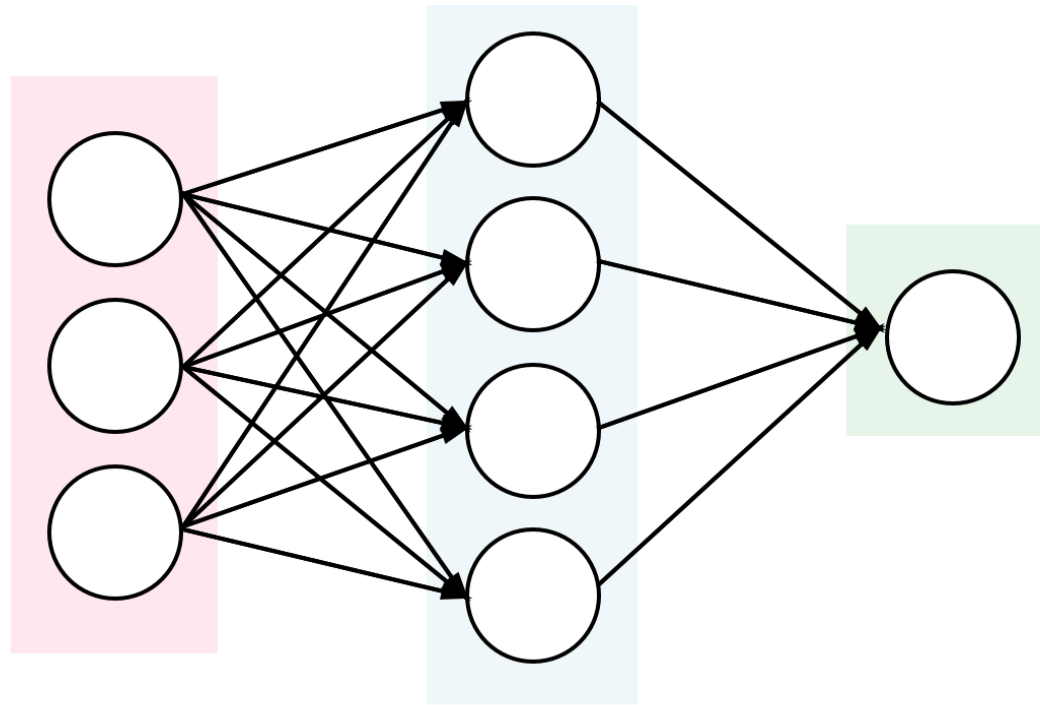
Constructing softmax regression

```
[14] model = keras.Sequential([  
    keras.layers.Flatten(input_shape=(28, 28)),  
    keras.layers.Dense(2, activation='softmax')  
])
```

```
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])
```

Artificial Neural Networks with one hidden layer

- Output: $\hat{\mathbf{y}} = \sigma(W_2 \text{relu}(W_1 \mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2)$



Constructing ANN

```
model = keras.Sequential([  
    keras.layers.Flatten(input_shape=(28, 28)),  
    keras.layers.Dense(50, activation='relu'),  
    keras.layers.Dense(10, activation='softmax')  
])
```

```
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])
```

Check model with model.summary()

```
model.summary()
```

```
Model: "sequential_4"
```

Layer (type)	Output Shape	Param #
=====		
flatten_4 (Flatten)	(None, 784)	0

dense_6 (Dense)	(None, 50)	39250

dense_7 (Dense)	(None, 10)	510
=====		

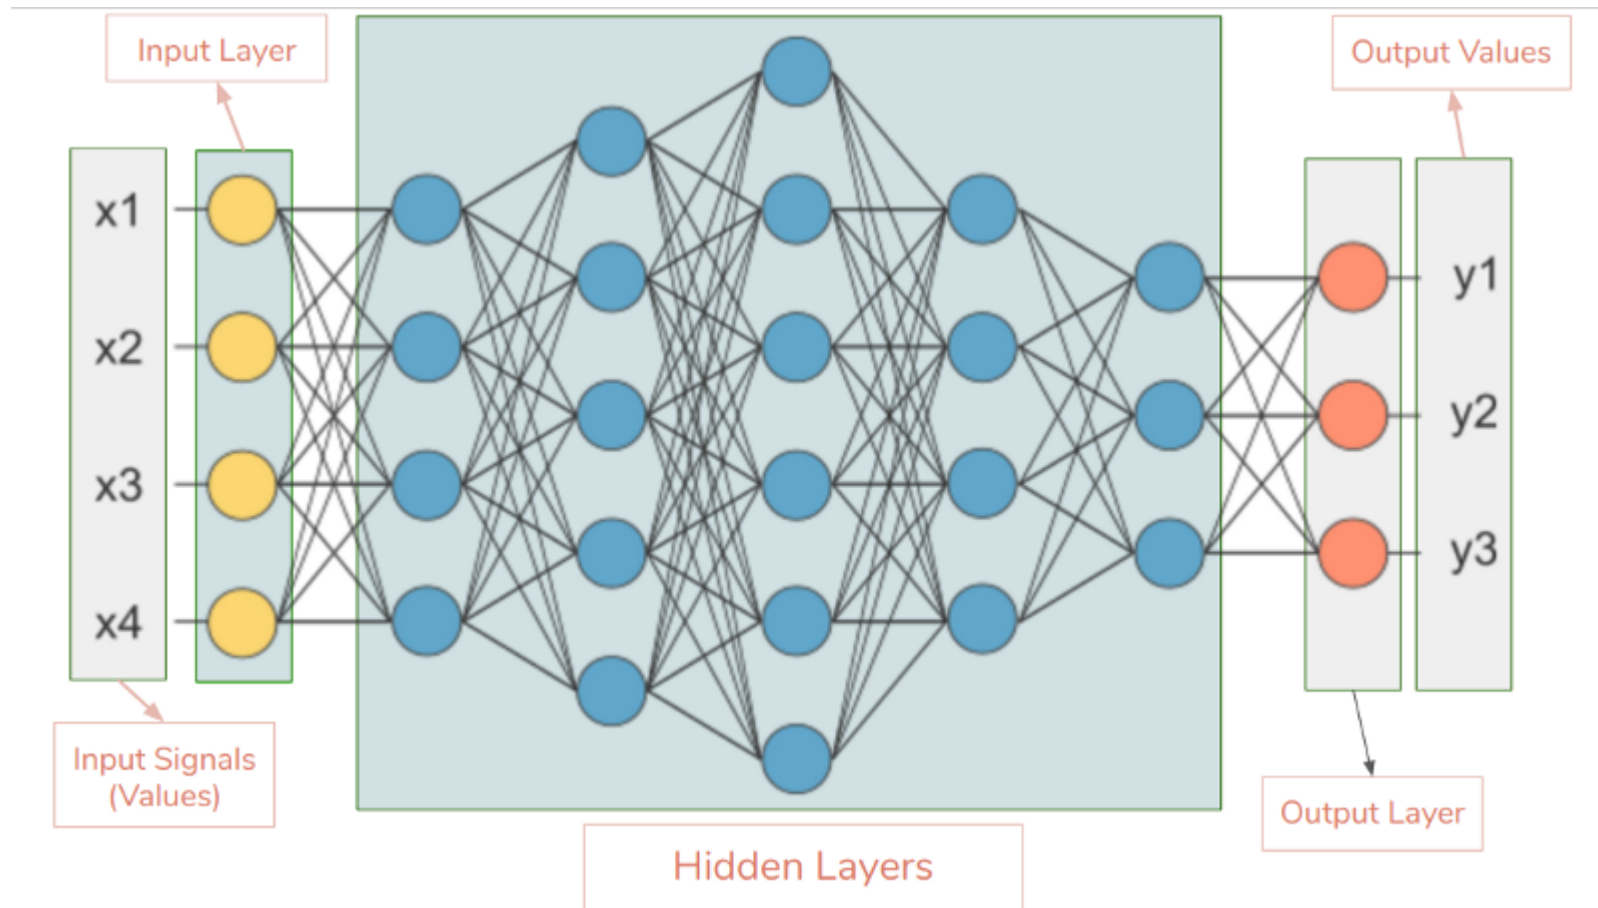
```
Total params: 39,760
```

```
Trainable params: 39,760
```

```
Non-trainable params: 0
```

ANN with multiple hidden layers

- **Ex)** $y = f_1(f_2(f_3(f_4(f_5(x))))))$

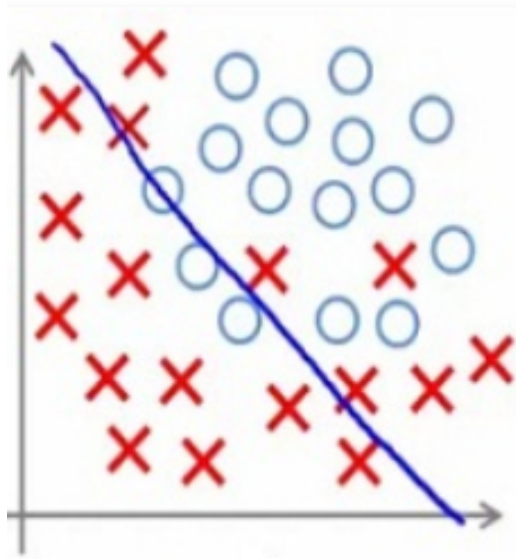


Practice

Train / Dev / Test set

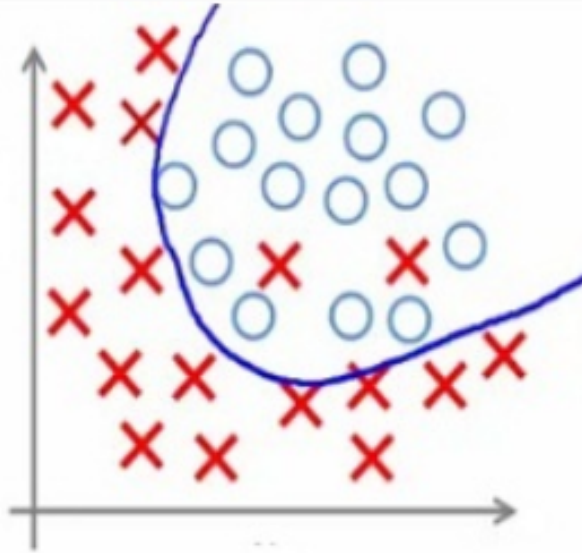
- Traditionally, 7:3 for train and dev or 6:2:2
- Or, 6:2:2
- With big data, 98:1:1 or use even larger train set
- It is important to use independent, separate Dev, Test set with different configuration for real evaluation

Overfitting / underfitting

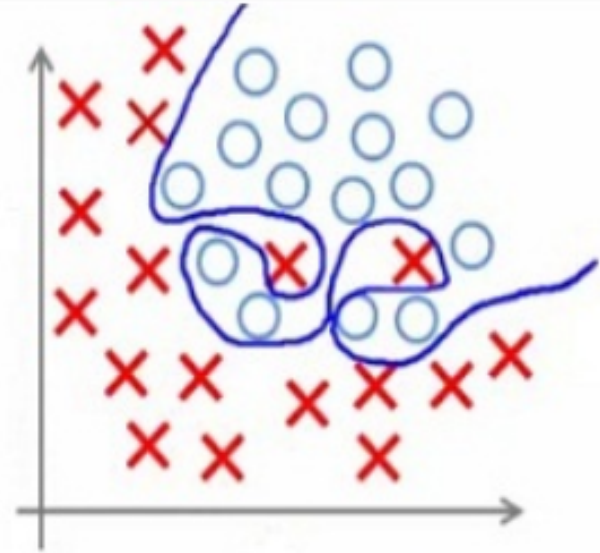


Under-fitting

(too simple to
explain the
variance)



Appropriate-fitting



Over-fitting

(forcefitting -- too
good to be true)

Check accuracy on train / dev

- Too good on train and low on dev imply overfitting
- Try to minimize accuracy(or EER, AUC, F1) on dev
- Check whether you have balanced or unbalanced data (train and dev, consider cost-sensitive learning when unbalanced)

Regularization

- Consider regularization when overfitted
- L1, L2

Ex) in logistic regression:

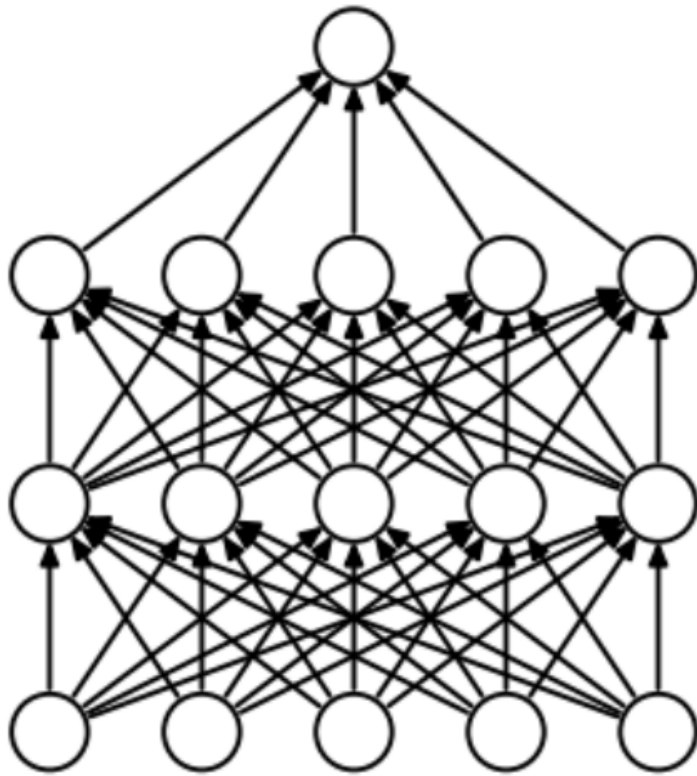
$$\mathbf{L1:} \quad C(\mathbf{w}, b) = \sum_{i=1}^m L(\hat{y}_i, y_i) + \frac{\lambda}{2m} \|\mathbf{w}\|^2 \quad \|\mathbf{w}\|^2 = \sum_{j=1}^{n_x} w_j^2$$

$$\mathbf{L2:} \quad C(\mathbf{w}, b) = \sum_{i=1}^m L(\hat{y}_i, y_i) + \frac{\lambda}{2m} \|\mathbf{w}\| \quad \|\mathbf{w}\| = \sum_{j=1}^{n_x} |w_j|$$

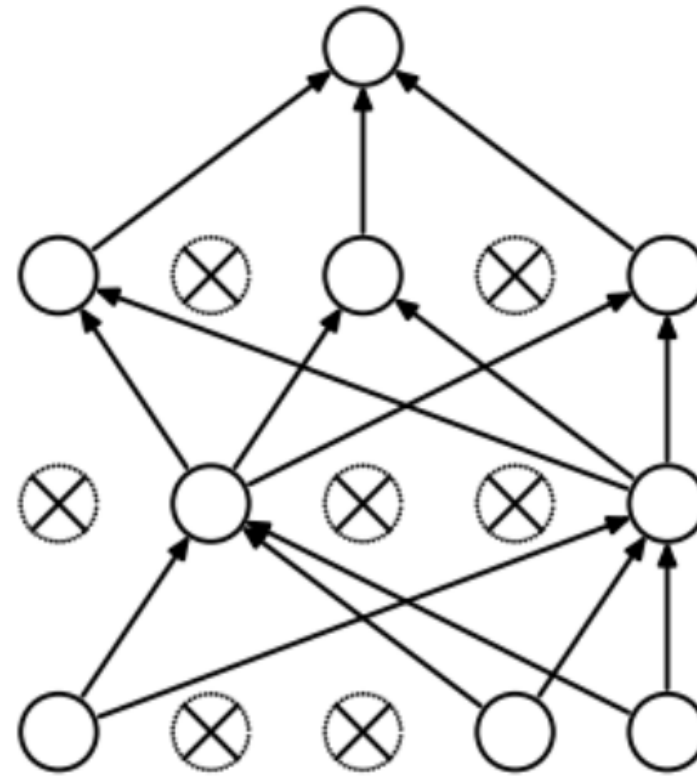
- Dropout

Dropout Regularization

- Reduce high dependency on few nodes (act like random forest)



(a) Standard Neural Net



(b) After applying dropout.

Codes for Regularization

- L2 :

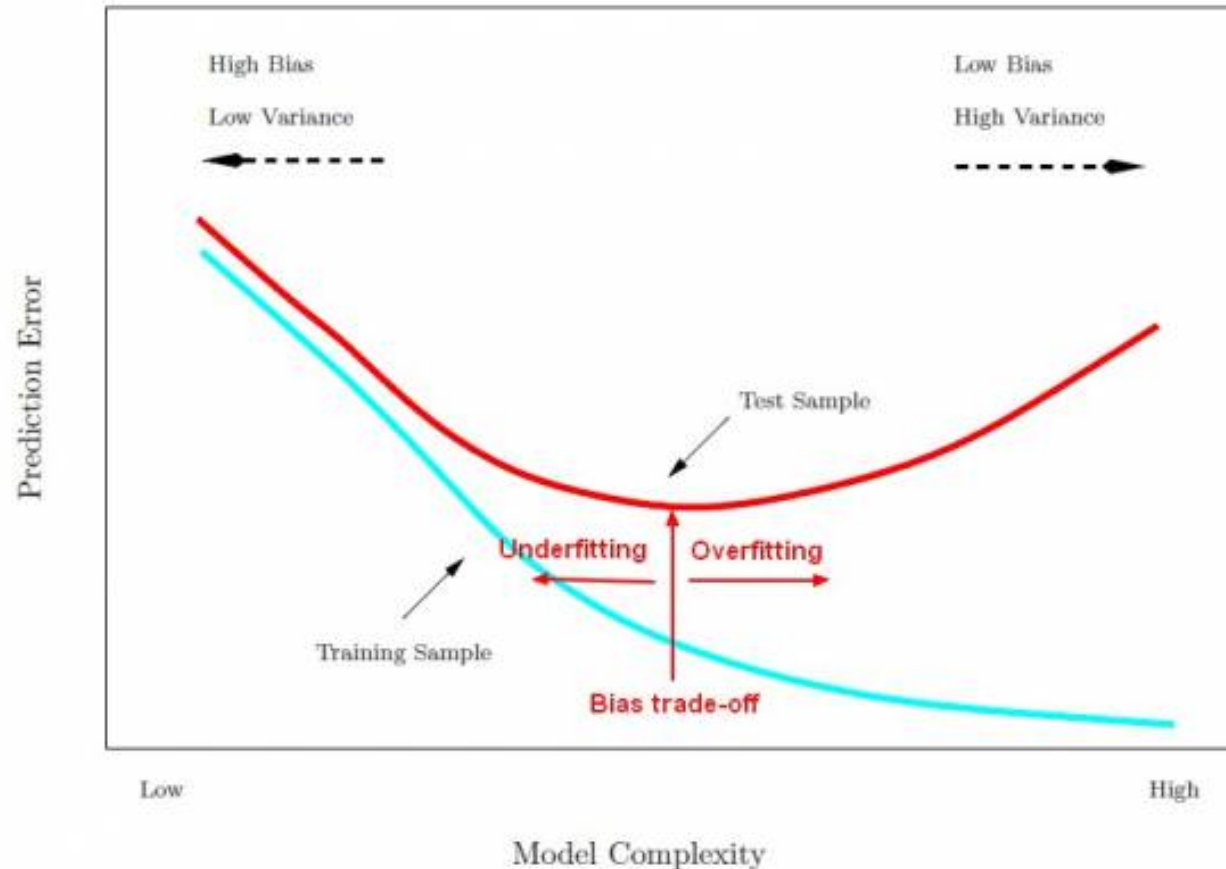
```
l2_model = keras.models.Sequential([
    keras.layers.Dense(16, kernel_regularizer=keras.regularizers.l2(0.001),
        activation='relu', input_shape=(NUM_WORDS,)),
    keras.layers.Dense(16, kernel_regularizer=keras.regularizers.l2(0.001),
        activation='relu'),
    keras.layers.Dense(1, activation='sigmoid')
])
```

- Dropout :

```
dpt_model = keras.models.Sequential([
    keras.layers.Dense(16, activation='relu', input_shape=(NUM_WORDS,)),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(16, activation='relu'),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(1, activation='sigmoid')
])
```

Early Stopping

- Reduce high dependency on few nodes (act like random forest)



Codes for Early Stopping

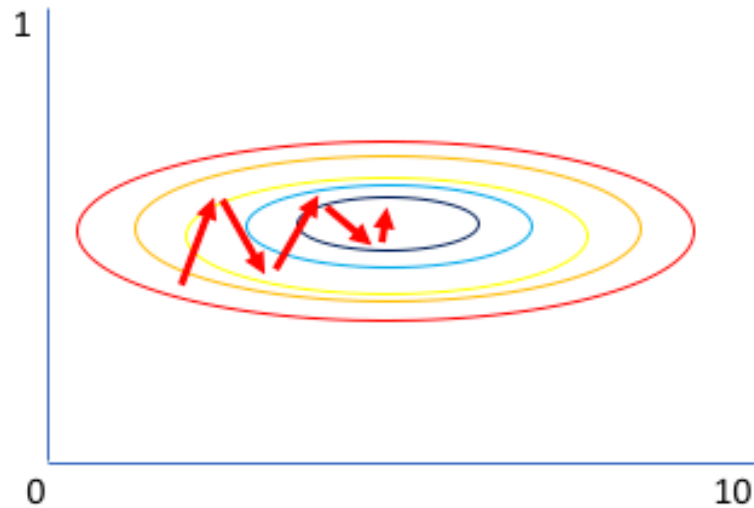
```
from keras.callbacks import EarlyStopping  
earlystop= EarlyStopping(monitor='val_acc', patience=3)
```

- **monitor** denotes the quantity that needs to be monitored and 'val_err' denotes the validation error.
- **Patience** denotes the number of epochs with no further improvement after which the training will be stopped.

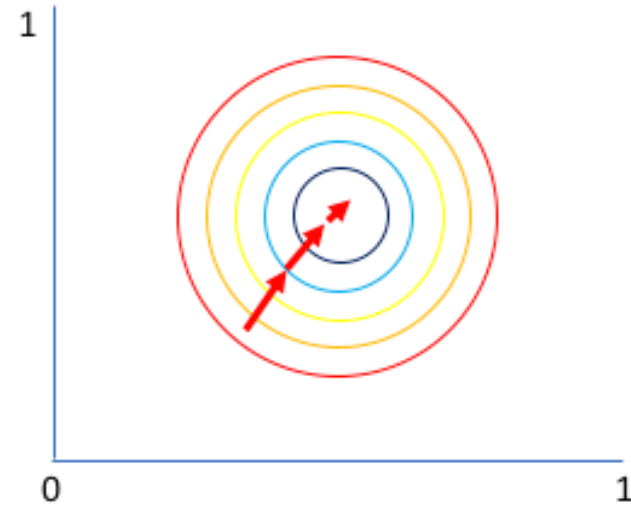
```
model.fit(train_images, train_labels, epochs=10,  
          validation_data = (test_images, test_labels),  
          callbacks=[earlystop] )
```

Normalizing inputs

Why normalize?



Gradient of larger parameter
dominates the update

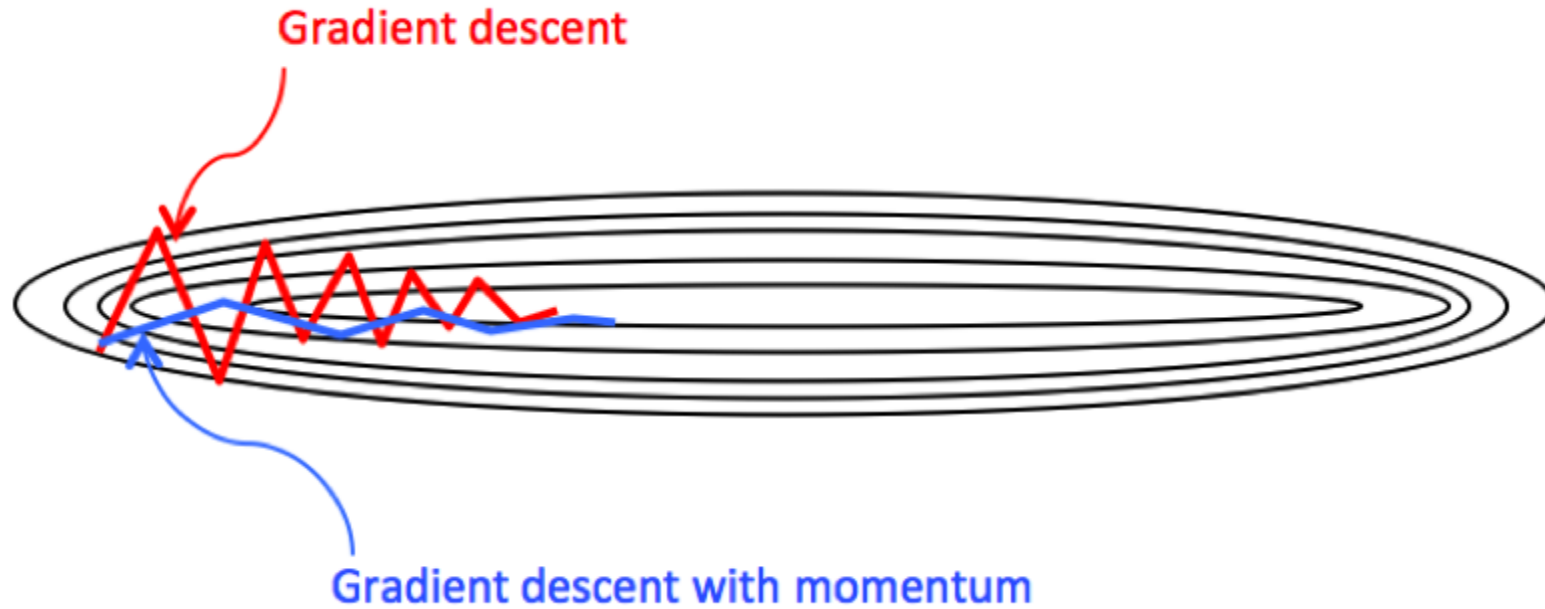


Both parameters can be
updated in equal proportions

Practice

<https://www.tensorflow.org/tutorials/keras/classification>

Gradient Decent with Momentum



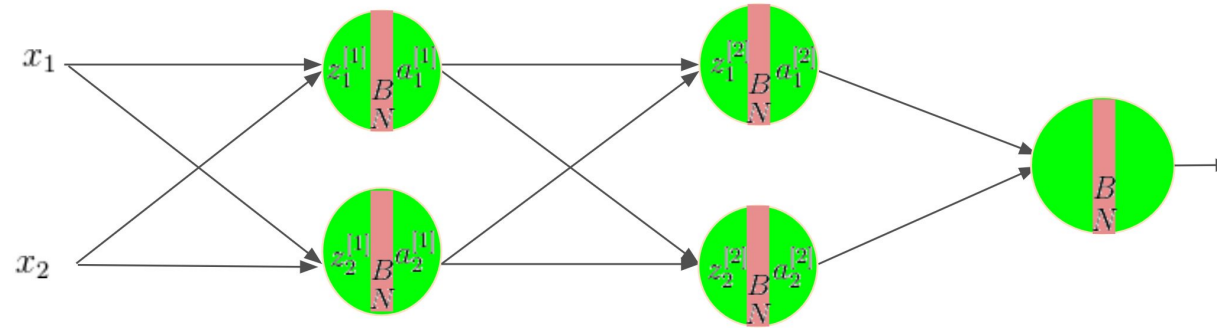
- RMSprop, and Adam optimizer

Learning rate decay / adaptive learning rate

- It is often useful to reduce learning rate as the training progresses
- Use **learning rate schedules** or **adaptive learning rate methods**

Batch Normalization

- Normalizing inputs to speed up learning



$$z^{[l]} = W^{[l]}a^{[l-1]} \longrightarrow \begin{aligned} \mu^{[l]} &= \frac{1}{m} \sum_i z^{[l](i)} \\ \sigma^{[l]2} &= \frac{1}{m} \sum_i (z^{[l](i)} - \mu^{[l]})^2 \\ z^{[l](i)}_{norm} &= \frac{z^{[l](i)} - \mu^{[l]}}{\sqrt{\sigma^{[l]2} + \epsilon}} \\ \tilde{z}^{[l](i)} &= \gamma^{[l]} z^{[l](i)}_{norm} + \beta^{[l]} \end{aligned} \longrightarrow a^{[l]} = g^{[l]}(\tilde{z}^{[l]})$$

- Use `keras.layer.BatchNormalization()`

Practice

<https://www.tensorflow.org/tutorials/keras/classification>

Thank you!
Q & A